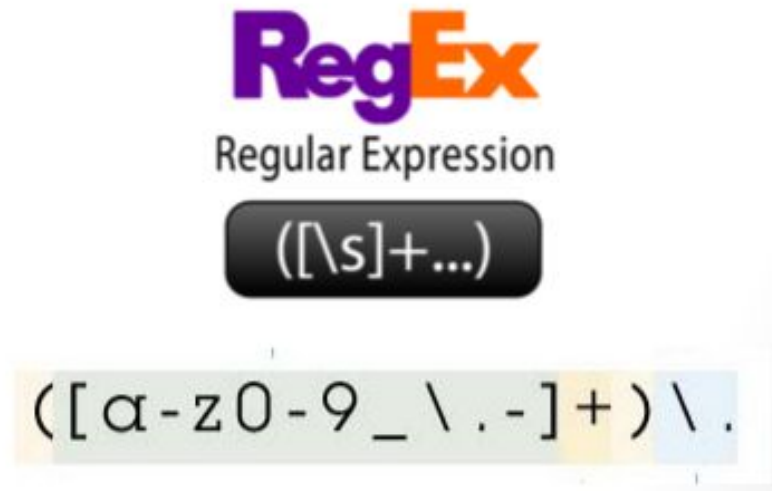


Topic 2 (Pt 2): Python Regular Expressions for NLP



Python RE

- A *regular expression* is a special sequence of characters that match or find other strings or sets of strings, using a specialized syntax held in a pattern.
- The Python **re** (called REs, or regexes, or regex patterns) is a tiny but highly specialized programming language embedded inside Python that provides full support for Perl-like regular expressions, made available through the re module.

Python RE (cont)

- Can be used with specified rules for the set of possible strings that you want to match; (ie., English sentences, e-mail addresses, html addresses, etc..).
- Ask questions such as:
 - “Does this string match the pattern?”, **OR**
 - “Is there a match for the pattern anywhere in this string?”
- Can also be used to modify/replace/split strings
 - <https://docs.python.org/3/howto/regex.html>
 - <https://docs.python.org/3/library/re.html>

Python RE (cont)

- There are various characters, which would have special meaning when they are used in regular expression.
- To avoid any confusion while dealing with regular expressions, we would use Raw Strings as `r'expression'` (the letter 'r', followed by the intended expression).

Common RE in Perl/Python, Unix Grep

RE	Example Patterns Matched
/woodchucks/	"interesting links to <u>woodchucks</u> and lemurs"
/a/	"Mary Ann stopped by Mona's"
/Claire_says,/	" "Dagmar, my gift please," <u>Claire</u> says,"
/DOROTHY/	"SURRENDER <u>DOROTHY</u> "
/!/	"You've left the burglar behind again!" said Nori

RE	Match	Example Patterns
/[wW]oodchuck/	Woodchuck or woodchuck	" <u>Woodchuck</u> "
/[abc]/	'a', 'b', or 'c'	"In uomini, in soldati"
/[1234567890]/	any digit	"plenty of <u>7</u> to 5"



A set of digits but match
one character at a time,
not the whole sequence



/[bc]oil/
matches boil OR coil and
NOT bcoil

Common RE in Perl/Python, Unix Grep

RE	Match
*	zero or more occurrences of the previous char or expression
+	one or more occurrences of the previous char or expression
?	exactly zero or one occurrence of the previous char or expression
{ <i>n</i> }	<i>n</i> occurrences of the previous char or expression
{ <i>n</i> , <i>m</i> }	from <i>n</i> to <i>m</i> occurrences of the previous char or expression
{ <i>n</i> , }	at least <i>n</i> occurrences of the previous char or expression

RE in Python

- Matching characters in Python

metachar	Meaning	Example
.	Matches any char except for newline	beg.n matches any character between 'beg' and 'n'. E.g., begun, begin, beg'n, etc
^	Complement or matches start of string	^z will match beginning of string with 'z' [^z] will match any char except z
\$	Matches end of string or before newline	'ther\$' will match any word ending with '...ther'
*	Matches 0 or more repetition	xy* matches x, xy or x followed by any no. of y's
+	Matches 1 or more repetition	ab+ matches ab and a followed by at least one b (e.g., abb, abbb, etc...)
?	Matches 0 or 1 repetitions of the preceding RE	ab? will match either 'a' or 'ab'.

RE in Python

metachar	Meaning	Example
{ }	Match exactly the no. of copies stated in { }	m{6} will match exactly 6 'm' characters (i.e, 'mmmmmm')
[]	Matches a set of character class	[abc] or [a-c] will match either a,b, or c [a-z] will match either a thru z (case sensitive)
\	Escape metacharacters	\[will match '[' in a pattern. E.g: [aeiou\], will match a,e,i,o,u or [
	Match either expression on each side	A B will match either A or B
()	Matches whatever regular expression inside parentheses	(abcdef) will match 'abcdef' (the whole sequence in parentheses)

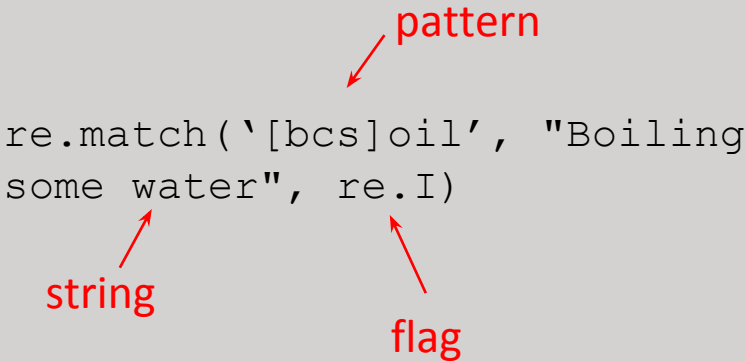
Advanced Operators

metachar	Expansion	Match	Example
\d	[0-9]	Any digit	There are 5 cats
\D	[^0-9]	Any non-digit	red ball
\w	[a-zA-Z0-9]	Any alphanumeric or space	RM9
\W	[^\w]	A non-alphanumeric	####
\s	[\r\t\n\f]	Whitespace, space or tab	
\S	[^\s]	Non whitespace	in school

Python re.match()

- This function attempts to match RE *pattern* to a *string* with optional *flags*.
- Syntax for match() with **2 required** and **1 optional** arguments/parameters:

`re.match(pattern, string, flags = 0)`

Param	Description	Example
pattern	Regular expression or pattern to be matched	 <pre>re.match('[bcs]oil', "Boiling some water", re.I)</pre> <p>The diagram shows the function call <code>re.match('[bcs]oil', "Boiling some water", re.I)</code>. Red arrows point from labels to the arguments: 'pattern' points to <code>'[bcs]oil'</code>, 'string' points to <code>"Boiling some water"</code>, and 'flag' points to <code>re.I</code>.</p>
string	String to be searched for matching with pattern	
flags	Indicate different flags using bitwise OR (<code> </code>). These are modifiers (explained in Table 3)	

RE modifiers: Option flags (Table 3)

Modifier	Description
re.I	Performs case-insensitive matching.
re.L	Interprets words according to the current locale. This interpretation affects the alphabetic group (<code>\w</code> and <code>\W</code>), as well as word boundary behavior (<code>\b</code> and <code>\B</code>).
re.M	Makes <code>\$</code> match the end of a line (not just the end of the string) and makes <code>^</code> match the start of any line (not just the start of the string).
re.S	Makes a period (dot) match any character, including a newline.
re.U	Interprets letters according to the Unicode character set. This flag affects the behavior of <code>\w</code> , <code>\W</code> , <code>\b</code> , <code>\B</code> .
re.X	Permits "cuter" regular expression syntax. It ignores whitespace (except inside a set <code>[]</code> or when escaped by a backslash) and treats unescaped <code>#</code> as a comment marker.

Python re.match()

- The *re.match* function returns a **match** object (single match) on success, **None** on failure.
- Use *group(num)* or *groups()* function of **match** object to get matched expression

```
>>> import re
>>> sent = "Boiling water is dangerous"
>>> re.match('[bcs]oil', sent, re.I)
<_sre.SRE_Match object; span=(0, 4), match='Boil'>

>>> mat = re.match('[bcs]oil', sent, re.I)
>>> mat.group()
'Boil'
```

returns an object

case
insensitive

returns a matched expression

RE : Matched subgroups

Match Object Methods	Description
group(num=0)	This method returns entire match (or specific subgroup num)
groups()	This method returns all matching subgroups in a tuple (empty if there weren't any)

```
>>> import re
>>> text = "A scripting language is more suitable for NLP than a compiled language"
>>> searchObj = re.search( r'(.*) is (.*) .*', text, re.M|re.I)
```

```
>>> if searchObj:
    print("searchObj.group() : ", searchObj.group())
    print("searchObj.group(1) : ", searchObj.group(1))
    print("searchObj.group(2) : ", searchObj.group(2))
else:
    print("Found nothing!!")
```

```
searchObj.group() : A scripting language is more suitable for NLP than a compiled language
searchObj.group(1) : A scripting language
searchObj.group(2) : more
```

Python re.match()

- More examples of matched expressions :
 - match a string that begins with with **0 or exactly 1 string of any combination**, followed by the word “grand”, “granda”, “grande”, “grandi”, “grando” or “grandu”:

```
>>> french2 = "Dans une grande ame tout est grand" means In a great mind  
everything is great"  
>>> re.match( '(.*?) grand[aeiou]*', french2, re.M|re.I).group()  
"Dans une grande"
```

- match the string that starts with **0 or more strings of any combination**, followed by the word “grand”, “granda”, “grande”, “grandi”, “grando” or “grandu”:

```
>>> re.match( '(.*) grand[aeiou]*', french2, re.M|re.I).group()  
"Dans une grande ame tout est grand"
```


Python search()

- `re.search()` searches for **first occurrence of RE *pattern*** within *string* with optional *flags*.
- Returns a **match** object on success, **None** on failure and uses *group(num)* or *groups()* to get matched expression
- Syntax : `re.search(pattern, string, flags = 0)`

```
>>> import re
>>> sent = "She is boiling some water and digging the soil"
>>> re.search('[bcs]oil', sent, re.I)
<_sre.SRE_Match object; span=(7, 11), match='boil'>
>>> srch = re.search('[bcs]oil', sent, re.I)
>>> srch.group()
'boil'
```

RE match() vs search()

- **re.match()** function only checks if the **RE** matches at the **beginning of the string**
- **match()** only reports a successful match which starts at position 0;

```
>>> print(re.match('gram', 'grammar'))  
<_sre.SRE_Match object at 0x0000000002F9FED0>  
>>> print(re.match('gram', 'programmer'))  
None
```

- **re.search()** scans forward through the string for a match.

```
>>> print(re.search('gram', 'grammar'))  
<_sre.SRE_Match object at 0x0000000002F9FED0>  
>>> print(re.search('gram', 'programmer'))  
<_sre.SRE_Match object at 0x0000000002F9FED0>  
>>> print(re.search('gram', 'programmer').group(0))  
gram
```


Greedy vs Ungreedy

- Repeating a regular expression as in **a*** attempts to consume as much of the pattern as possible because of its **greedy** nature **.***.

```
>>> |s = '<html><head><title>Title</title>'  
>>> print(re.match('<.*>', s).group(0))  
<html><head><title>Title</title>
```

- Use the **non-greedy** qualifiers ***?**, **+?**, **??**, or **{m,n}?**, which match as little text as possible.

```
>>> print(re.match('<.*?>', s).group(0))  
<html>
```

Modifying Strings

Method/Attribute	Purpose
<code>split()</code>	Split the string into a list, splitting it wherever the RE matches
<code>sub()</code>	Find all substrings where the RE matches, and replace them with a different string
<code>subn()</code>	Does the same thing as <code>sub()</code> , but returns the new string and the number of replacements

- **.sub** Returns the string obtained by **replacing the leftmost non-overlapping occurrences** of the RE in *string* by the replacement *replacement*

Another way of defining a pattern is to compile it first.

```
>>> p = re.compile( '(blue|white|red)' )
>>> p.sub( 'colour', 'blue socks and red shoes' )
'colour socks and colour shoes'
>>> p.sub( 'colour', 'blue socks and red shoes', count=1 )
'colour socks and red shoes'
```

Modifying Strings

- The *.subn()* method does the same work, but returns a 2-tuple containing the new string value and the number of replacements performed

```
>>> p = re.compile( '(blue|white|red)')
>>> p.subn( 'colour', 'blue socks and red shoes')
('colour socks and colour shoes', 2)
>>> p.subn( 'colour', 'no colours at all')
('no colours at all', 0)
```

RE in Python

- More RE methods

Method/Attribute	Purpose
<code>match()</code>	Determine if the RE matches at the beginning of the string.
<code>search()</code>	Scan through a string, looking for any location where this RE matches.
<code>findall()</code>	Find all substrings where the RE matches, and returns them as a list.
<code>finditer()</code>	Find all substrings where the RE matches, and returns them as an <i>iterator</i> .

```
>>> import re
>>> m = re.search('(?<=abc)def', 'abcdef')
>>> m.group(0)
'def'
```

```
>>> m = re.search('(?<=-)\w+', 'spam-egg')
>>> m.group(0)
'egg'
```

Example problem for RE in NLP

- Write a regular expression that finds all the instances of the word 'the' in the text:

“The boy eats another cake in the kitchen”

`/the/` : This pattern only finds lower case 'the'

`/[tT]he/` : This pattern finds all words with 'the' or 'The' in it

`/\b[tT]he\b/` : This pattern finds all 'the' or 'The' (`\b` matches boundaries)

Example solution for RE in NLP

- Matching ‘the’ that precedes an underscore or numbers (e.g., “the_” or “the25”):

`/[^a-zA-Z][tT]he[^a-zA-Z]/`

- Problem: This pattern will not match *the* when it begins a line or ends a line (e.g., “*The book...*” or “*... the.*”)
- Solution:

`/(^|[^a-zA-Z])[tT]he([a-zA-Z]|$)/`

RE Errors

- In RE, the process involves fixing two kinds of errors:
 - Matching strings that we should not have matched (*there, then, other*)
 - False positives (Type I error)
 - Not matching strings that we should have matched (*The*):
 - False negatives (Type II error)

Errors

- Reducing error rates for an application involves 2 efforts:
 - Increase accuracy or precision (i.e., minimize false positives)
 - Increase coverage or recall (i.e., minimize false negatives)

